

SUMMARY

EXIST Gründerstipendium recipient supporting venture-backed AI startup development.

Experienced in deploying LLMs and generative models into scalable backend infrastructures, optimizing inference pipelines, and translating ambiguous client requirements into robust AI solutions. Experienced in architecting end-to-end AI infrastructure in early-stage startup environments.

PROFESSIONAL EXPERIENCE

Founding Engineer

Berlin, Germany

Kivo - Software Development and Process optimization

Jun 25 - Present

- Built production-ready LLM pipelines for structured documents (invoices, contracts, supplier quotes) using embedding models and vector databases
- Designed similarity scoring and ranking systems for real-time candidate-job matching using optimized embedding search.
- Led backend architecture (FastAPI, PostgreSQL, Docker, cloud deployment) ensuring reliability and low-latency inference.
- Developed evaluation frameworks to benchmark precision, recall, and system performance in production environments.
- Implemented scalable Python ETL pipelines handling large-scale document ingestion and preprocessing.

AI Development Intern

Berlin, Germany

Rainmaker Society- Career support and a professional candidate matching network

Jan 25 - May 25

- Built LLM-powered document intelligence pipeline converting unstructured CVs and job descriptions into structured embeddings.
- Designed ranking algorithms and similarity scoring logic for automated candidate matching systems.
- Implemented scalable backend processing pipelines for large-volume document handling.
- Defined LLM evaluation framework (benchmarking precision vs. recall) to improve matching precision and system performance.

Research Intern - Data Analytics and Modeling

Berlin, Germany

Atiko7 Climate Labs - Climate-tech digital product lab applying AI to sustainability

Apr 24 - Jun 24

- Improved DeepForest model accuracy by ~10% via dataset refinement and retraining on AWS SageMaker.
- Designed ETL pipelines for large-scale environmental datasets aligned to stakeholder KPIs.
- Fine-tuned generative models (Stable Diffusion) ensuring quality standards for production grade use cases.

ACADEMIC EXPERIENCE

Research Assistant - ESMT Berlin

Berlin, Germany

ESMT Berlin

Aug 24 - Sep 24

- Built Python pipelines to extract financial data from SEC EDGAR filings delivering clean, structured data for financial analytics.
- Designed SQL-compatible data models for analytics dashboards and KPI reporting.

Research Intern ML/AI

Mumbai, India

Indian Institute of Technology (BHU)

Aug 21 - Mar 22

- Developed a dlib model from scratch for cattle recognition and achieved over 90% ranked accuracy
- Deployed search efficiency enhancements using Kibana Dashboard Visualizations and Elastic Search filters

EDUCATION

Master of Science in Analytics and Artificial Intelligence

Sep 23 - Sep 25

ESMT Berlin, Germany

Master Thesis: Hybrid CNN-Transformer Models and Pretrained ViTs under resource constraints

Bachelor of Technology in Computer Science and Engineering

Jul 19 - Jul 23

VIT Vellore, India

TECHNICAL SKILLS

- **AI/ML Frameworks:** PyTorch, HuggingFace, LangChain, LLM APIs, Vector Databases, Embedding Pipelines
- **Backend and Infrastructure:** Python, FastAPI, Docker, PostgreSQL, Supabase, REST APIs
- **Cloud & Deployment:** AWS SageMaker, GCP Vertex AI, Azure Functions
- **Optimization/Analytics:** Ranking Algorithms, Similarity Scoring, Dynamic Pricing Models, KPI Modeling

LANGUAGES

English: Proficient, German: Intermediate